STAT 230A Final Project: Predicting BART Ridership using Hourly Destination Data

Drew Galloro & Benjamin Khothsombath

May 13, 2026

1 Introduction

Bay Area Rapid Transit (BART) serves most parts of the San Francisco Bay Area and is a vital service for locals and tourists traveling the area. A model predicting ridership from existing data would be instrumental for the agency, as understanding ridership patterns would help set schedules, justify new housing near stations, or guide service reductions.

In this paper, we engineer features that we believe are valuable in predicting hourly ridership. From these features, we construct numerous linear models and identify a Poisson GLM as the best model for prediction.

2 Data and Domain Problem

Our original data is sourced from BART's ridership reports. Specifically, the hourly trip datasets, which run from 2018 to 2025. The 2026 dataset has not been updated yet. An observation consists of a date, hour, origin station, destination station, and the number of riders that completed that trip within that hour.

Due to the massive number of rows that the dataset contains by sourcing every possible combination of stations for every hour over a span of 7 years, we perform two preprocessing steps to make the project simpler: 1) we reduce our training data to only that of 2023-2024, as we observed that ridership was much different after the COVID-19 pandemic. 2) We aggregated origin-destination station pairs to the level of only the destination. Doing this simplifies the problem from trying to predict the number of riders that took a specific trip during an hour to predicting the number of riders that exited a station during a given hour. Furthermore, it shrinks the size of the data from $\sim$ 6–7 GB to $\sim$ 0.1 MB, which is much more feasible to work with.

Prior to modeling, we engineered a few features.

1. `is_holiday`: an indicator variable for whether a given day is a holiday, using the `tis` package
2. `is_giants_home` and `is_as_home`: indicator variables for whether a home baseball game is being played by the SF Giants or Oakland A's, respectively
3. `warriors_at_chase`: an indicator variable for whether a home game is being played by the Golden State Warriors
4. `season`: a factor of 4 levels indicating the season of the year
5. `day`: a factor with 3 levels - `weekday/saturday/sunday`

It is also worth mentioning that we converted the `hour` variable to a factor. The reason for this is that ridership is a nonlinear function of hour, as ridership does not strictly increase as the day goes on. We

could have modeled the hour as a polynomial, but we opted to leave it as a factor to allow the models to better learn ridership at a particular hour.

3 Our Final Regression Model

Our final model is a Poisson GLM with the above features, as well as an interaction term between the hour and day. We performed model selection on a model trained on data from 2023-2024 and refit this model to include 2025 data after having shown it achieved the lowest RMSE on the 2025 hold-out set. Overall, the model includes 128 coefficients, all but one of which are highly significant at a Bonferonni-corrected significance level of 1%. The spread of coefficients for each variable is shown in Table 1.

Term	**# Coefficients**	**# Significant**	**Mean $\lvert\hat{\beta}\rvert$**
Destination Station	49	49	0.829
Hour	23	23	1.384
Day	2	2	0.335
Holiday	1	1	0.696
Giants Home Game	1	1	0.031
A's Home Game	1	1	0.008
Warriors Home Game	1	1	0.021
Season	3	3	0.077
Hour $\times$ Day	46	45	0.822
Total	**128**	**127**	

Table 1: Poisson GLM coefficient summary. Mean $|\hat{\beta}|$ reports the average absolute coefficient within each feature group on the log-rate scale, providing a rough measure of feature group importance. Significance assessed at $\alpha = 0.05$ with Bonferroni correction.

From this, we see that roughly, the most important variables for predicting BART ridership are the time of day, day of the week, and whether the day is a holiday. The sports indicators are interestingly, rather weak indicators of overall ridership, suggesting that perhaps a better model would consider the location and time of sports games rather than the overall effect of sports games on ridership. Practically, this would mean that overall BART schedules shouldn't adjust system-wide schedules to accommodate sports games, though a station-level analysis near Oracle Park or the Chase Center may be worth pursuing.

4 Discussion

Although the Poisson regression model was selected as the best model for predicting ridership, it comes with many flaws and nuances. One particular drawback of the Poisson GLM is that it assumes that the response is Poisson distributed and thus, has variance equal to its mean. This is a rather strong assumption which tends to not hold in practice, so we anticipated a negative binomial GLM performing better. However, we had chosen MSE on a test set as our metric of choice for evaluating models, so we opted to use this model. Furthermore, we performed the likelihood ratio test to test for overdispersion by comparing this model to a NB model and obtained a chi-squared value of 3.3×10^7, suggesting heavy overdispersion and model misspecification. However, model misspecification is not necessarily a problem for prediction, as we are simply trying to minimize the empirical expected risk, which is possible even if we didn't properly model the distribution correctly.

We also observed some classical model assumptions that are violated in this model. We discuss some of them further in Section 6, but the primary one that BART administration may be interested in is the heteroskedasticity. Heteroskedasticity is not typically a problem for GLMs, but there are some practical implications. In particular, certain areas where the residuals are abnormally high reflect "pockets" where the model is not able to properly predict ridership. This ties back to the misspecification problem, but in a different way: there are features missing that we did not account for. For example, one of them may be that

there are some days where protests or concerts occur and some stations experience large ridership, which we didn't account for. Another is that there may be times where the system is experiencing maintenance issues, forcing riders to exit at a particular station. As evidence of these "pockets" where the model doesn't account for some variables, we observe the following plots. Figure 1b shows that the data points with the highest residuals are the ones with smaller leverage, which corresponds to observations with few indicators. The rootogram in Figure 1a also shows that for small ridership values, the model underpredicts the true ridership while overpredicting when the true ridership is ∼100–250 for a given station-hour.

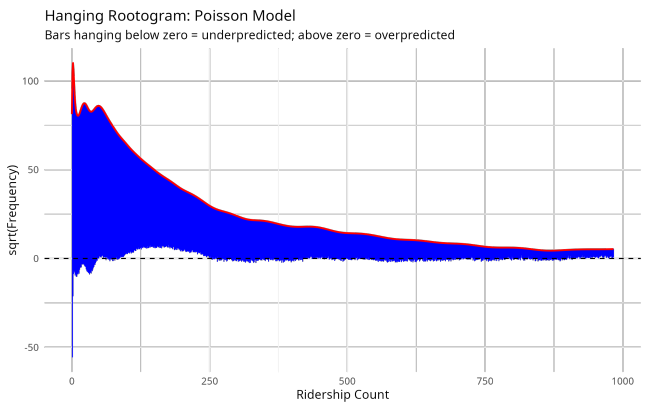

(a) Hanging rootogram of the final Poisson model. Bars below zero indicate underprediction; bars above zero indicate overprediction.

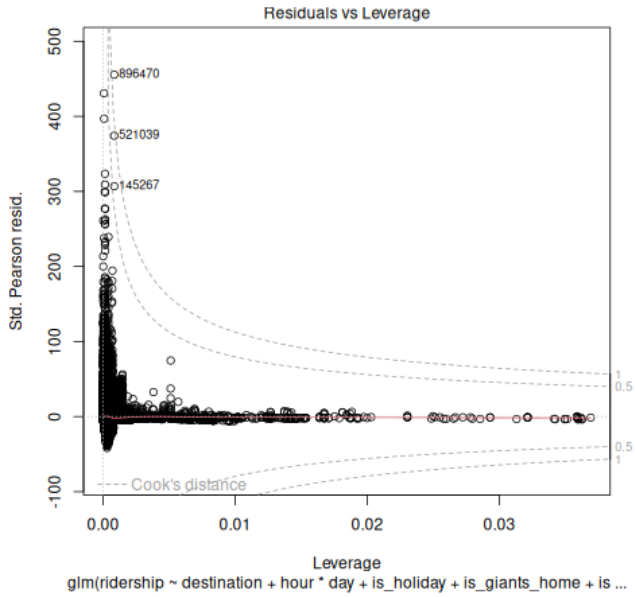

(b) Leverage vs. standardized residuals. High-residual points cluster at low leverage, suggesting missing predictors rather than influential observations.

Figure 1: Diagnostic plots for the final Poisson GLM fitted on the full dataset (2023–2025).

5 Conclusion

We conclude that while our model provides a strong baseline for predicting a destination's ridership, some further work is required before adopting this model for use in BART administrative decision making. Our model, when trained on 2023-2024 data and tested on 2025 data, achieved a RMSE of 174, meaning that we expect our estimate of the amount of riders exiting a given station to be off by roughly 174. We believe this amount of error is too high for us to be comfortable using this information for BART planning and that further work will need to be done.

We acknowledge that all of the features we've engineered so far are statistically significant, but further feature engineering may identify gaps in our model that would help predict ridership for a given station. To start, we recommend further inspecting areas/times where the model deviates most from reality and building features to account for this. As an example, our penultimate model identified the hour × day interaction as a potential gap, so we incorporated it into our final model. After doing this, we also identified some high-residual points clustered around Coliseum on the weekends, which suggests adding a Coliseum × sports-game or Coliseum × day interaction to the existing model as a first step for this iterative process.

6 Additional Work

6.1 Model Selection

For our model selection methodology, we pre-registered the following framework: we'll begin by first conducting feature engineering and using the raw features we've derived. Then, we'll construct the following models based on data from 2023-2024, and select the best model based on whichever achieves the lowest test MSE on the 2025 data. We selected the Basic OLS as a baseline model and opted for various models that account for the right-skewed nature of the ridership. The results are shown below.

Model	**Test MSE**	**Test RMSE**
Basic OLS LM	41,677	204
Log-Transformed LM (bias-corrected)	47,682	218
Poisson GLM	**32,969**	**182**
Negative Binomial GLM	46,721	216
Ridge Poisson ($\lambda = 7.082$)	33,584	183
Ridge NB ($\lambda = 0.599$)	43,978	210
Zero-Inflated NB	46,697	216

Table 2: Test set performance of all preliminary models evaluated. MSE and RMSE are reported on the original ridership scale (riders per station-hour). The Poisson GLM achieves the lowest test MSE and is selected as the final model.

For Model 2, the predicted value of the log-transformed model is on the log-scale. However, exponentiating the predicted value will give a biased estimate of the predicted value. Thus, we perform a correction by multiplying by $\exp(\hat{\sigma}^2/2)$ to correct for this bias, though this assumes normality of residuals.

Regarding the selection of λ for the ridge regressions, we ran a 5-fold CV using time-based splits for the Poisson Ridge (Model 5). In doing so, we'd technically be using future data (for example, fold 4) to predict past data (for example, fold 3), but we're not currently concerned with autocorrelation. For the NB Ridge (Model 6), we encountered memory issues when trying to perform this procedure, so we instead selected λ using a 2023 as our inner-training set and 2024 as our validation set and refit the model on the full training data after selecting λ.

As previously stated, the Poisson model outperforming the negative binomial model went against our expectations considering the overdispersion. However, this isn't too surprising when considering that the NB model also requires estimating θ to better model the variance. Although the Poisson model does not correctly capture the variance of the distribution of ridership, it seems that it's still able to predict the conditional mean better.

After performing this analysis, we did a deep dive into where our first Poisson model performed worst and identified weekday peak hours as performing the worst in the model. Table 3 shows that our first Poisson model was unable to identify differences in rush-hour traffic for weekdays vs weekends, which motivated our introduction of the interaction term for the final model. This model achieved a test MSE of 30,254 and test RMSE of 174 prior to being refit on the entire dataset.

6.2 Model Assumptions

6.2.1 Autocorrelation

Since we are looking at ridership over time, we naturally acknowledge that this problem may be better suited to time-series methods. Naturally, we expect to see seasonality in ridership and dependence between ridership at similar times, which the plots in Figure 2 demonstrate. We attempt to account for this by accounting for the `day`, `hour`, and `season`, but a time-series analysis is likely more robust.

Day Type	Period	RMSE	MAE	n
Weekday	Off-Peak	84.77	39.40	197,085
Weekend	Off-Peak	99.59	41.21	71,202
Weekday	PM Peak	240.83	157.39	39,109
Weekday	AM Peak	444.74	170.29	38,658
Weekend	PM Peak	111.55	49.33	15,575
Weekend	AM Peak	107.93	74.86	14,465

Table 3: Stratified prediction performance of the Poisson GLM (trained on 2023-2024 data, no interaction) by day type and time period. AM Peak covers hours 7–9, PM Peak covers hours 16–18, and Off-Peak covers all remaining hours. RMSE and MAE are on the original ridership scale (riders per station-hour).

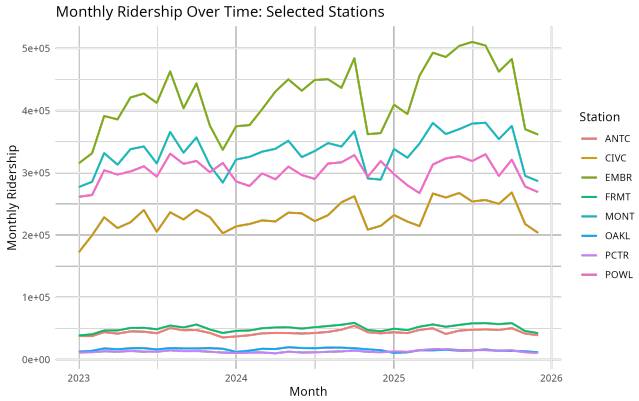

(a) Monthly ridership for selected stations. We try to incorporate the overall seasonality by accounting for the season, but our approach likely underfits the true effect.

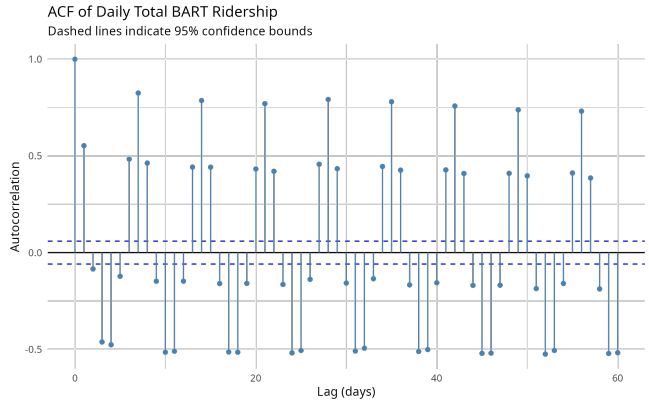

(b) ACF of overall ridership over days. The highly correlated time lags, especially at multiples of 7 days further demonstrates autocorrelation.

Figure 2: Plots demonstrating that Time Series Analysis may be more appropriate.

6.2.2 Multicollinearity

We demonstrate multicollinearity via Table 4. The $\text{GVIF}^{1/2\text{df}}$ for `hour`, `day`, and their interaction demonstrate the high collinearity, which is a classical violation of assumptions, but does not present a problem for prediction since future data will reflect the same correlation structure for these variables.

Term	GVIF	Df	$\text{GVIF}^{1/(2\cdot\text{Df})}$
Destination	1.00×10^{0}	49	1.000
Hour	8.40×10^{24}	23	3.482
Day	3.62×10^{3}	2	7.758
Holiday	1.01×10^{0}	1	1.004
Giants Home Game	1.21×10^{0}	1	1.099
A's Home Game	1.17×10^{0}	1	1.084
Warriors Home Game	1.07×10^{0}	1	1.034
Season	1.47×10^{0}	3	1.067
Hour × Day	6.60×10^{27}	46	2.006

Table 4: Generalized Variance Inflation Factors (GVIF) for the final Poisson GLM. For terms with more than one degree of freedom, $\text{GVIF}^{1/(2\,\cdot\,\text{Df})}$ is the appropriate analog to the square root of VIF and is comparable across terms. The large GVIF values for `hour`, `day`, and their interaction reflect expected multicollinearity introduced by the interaction term, which is not a concern for prediction.

7 Statement of Contribution

- Drew Galloro worked on the data pulling scripts, data cleaning, exploratory data analysis, and feature engineering.
- Benjamin Khothsombath worked on the feature engineering, modeling, and model diagnostics.

8 Appendix

High-Residual Observations

Destination	Hour	Day	Holiday	Fitted	Pearson Resid.	Leverage	Abs. Resid.
COLS	12	Sunday	FALSE	52.67	435.87	3.55×10^{-5}	435.87
COLS	12	Sunday	FALSE	52.67	416.72	3.55×10^{-5}	416.72
COLS	12	Sunday	FALSE	52.67	347.27	3.55×10^{-5}	347.27
COLS	15	Saturday	FALSE	96.03	327.47	5.82×10^{-5}	327.47
COLS	12	Sunday	FALSE	54.38	313.34	3.70×10^{-5}	313.34
COLS	11	Sunday	FALSE	51.63	300.66	3.49×10^{-5}	300.66
LAKE	12	Saturday	FALSE	75.46	289.46	4.53×10^{-5}	289.46
COLS	15	Saturday	FALSE	96.03	286.85	5.82×10^{-5}	286.85
COLS	12	Sunday	FALSE	52.67	252.75	3.55×10^{-5}	252.75
16TH	11	Saturday	FALSE	142.71	232.48	6.13×10^{-5}	232.48
MLBR	1	Weekday	TRUE	1.72	222.19	8.67×10^{-6}	222.19
COLS	12	Sunday	FALSE	52.67	211.27	3.55×10^{-5}	211.27
LAKE	12	Saturday	FALSE	77.91	209.15	4.56×10^{-5}	209.15
COLS	12	Sunday	FALSE	52.67	197.77	3.55×10^{-5}	197.77
12TH	12	Saturday	FALSE	120.66	192.12	5.63×10^{-5}	192.12
LAKE	11	Saturday	FALSE	76.37	189.45	4.49×10^{-5}	189.45
LAKE	13	Saturday	FALSE	78.19	186.35	4.65×10^{-5}	186.35
MLBR	1	Weekday	TRUE	1.72	181.76	8.67×10^{-6}	181.76
COLS	19	Saturday	FALSE	78.25	180.74	4.83×10^{-5}	180.74
COLS	13	Sunday	FALSE	54.57	176.98	3.65×10^{-5}	176.98

Table 5: Top 20 high-residual observations with below-median leverage from the penultimate Poisson GLM. The concentration of COLS (Coliseum) on weekends suggests an unmodeled effect, such as a Coliseum $\times$ day or Coliseum $\times$ sports game interaction.

Full Coefficient Table

Term	Estimate	Std. Error	z value	Pr($>\|z\|$)
(Intercept)	3.9522	0.0020	1987.56	$< 2 \times 10^{-16}$
Destination (reference: 12TH)				
16TH	0.1922	0.0006	299.98	$< 2 \times 10^{-16}$
19TH	-0.0346	0.0007	-51.17	$< 2 \times 10^{-16}$
24TH	0.1513	0.0006	234.01	$< 2 \times 10^{-16}$
ANTC	-1.0238	0.0009	-1106.16	$< 2 \times 10^{-16}$
ASHB	-0.8994	0.0009	-1019.49	$< 2 \times 10^{-16}$
BALB	-0.2102	0.0007	-296.47	$< 2 \times 10^{-16}$
BAYF	-0.6287	0.0008	-781.73	$< 2 \times 10^{-16}$
BERY	-1.1321	0.0010	-1178.46	$< 2 \times 10^{-16}$
CAST	-1.3903	0.0011	-1307.46	$< 2 \times 10^{-16}$
CIVC	0.6189	0.0006	1052.12	$< 2 \times 10^{-16}$
COLM	-1.1146	0.0010	-1167.91	$< 2 \times 10^{-16}$
COLS	-0.6051	0.0008	-758.24	$< 2 \times 10^{-16}$
CONC	-0.6560	0.0008	-808.54	$< 2 \times 10^{-16}$
DALY	-0.0282	0.0007	-41.69	$< 2 \times 10^{-16}$
DBRK	0.1795	0.0006	279.35	$< 2 \times 10^{-16}$
DELN	-0.2575	0.0007	-358.54	$< 2 \times 10^{-16}$
DUBL	-0.4797	0.0008	-623.66	$< 2 \times 10^{-16}$
EMBR	1.2248	0.0005	2270.21	$< 2 \times 10^{-16}$
FRMT	-0.8920	0.0009	-1013.40	$< 2 \times 10^{-16}$
FTVL	-0.2195	0.0007	-308.74	$< 2 \times 10^{-16}$
GLEN	-0.4301	0.0008	-569.27	$< 2 \times 10^{-16}$
HAYW	-0.8022	0.0009	-940.73	$< 2 \times 10^{-16}$
LAFY	-1.1422	0.0010	-1184.31	$< 2 \times 10^{-16}$
LAKE	-0.4659	0.0008	-609.80	$< 2 \times 10^{-16}$
MCAR	-0.2678	0.0007	-371.60	$< 2 \times 10^{-16}$
MLBR	-0.8806	0.0009	-1002.63	$< 2 \times 10^{-16}$
MLPT	-1.3431	0.0010	-1287.78	$< 2 \times 10^{-16}$
MONT	0.9936	0.0006	1789.51	$< 2 \times 10^{-16}$
NBRK	-1.0082	0.0009	-1098.94	$< 2 \times 10^{-16}$
NCON	-2.0760	0.0014	-1460.39	$< 2 \times 10^{-16}$
OAKL	-2.0323	0.0014	-1447.50	$< 2 \times 10^{-16}$
ORIN	-1.3910	0.0011	-1307.17	$< 2 \times 10^{-16}$
PCTR	-2.2464	0.0015	-1459.58	$< 2 \times 10^{-16}$
PHIL	-0.7337	0.0008	-881.05	$< 2 \times 10^{-16}$
PITT	-0.8853	0.0009	-1008.14	$< 2 \times 10^{-16}$
PLZA	-0.8437	0.0009	-975.44	$< 2 \times 10^{-16}$
POWL	0.8918	0.0006	1583.57	$< 2 \times 10^{-16}$
RICH	-0.7302	0.0008	-875.64	$< 2 \times 10^{-16}$
ROCK	-0.6359	0.0008	-788.73	$< 2 \times 10^{-16}$
SANL	-0.4138	0.0008	-550.40	$< 2 \times 10^{-16}$
SBRN	-1.1686	0.0010	-1199.50	$< 2 \times 10^{-16}$
SFIA	-0.0910	0.0007	-132.52	$< 2 \times 10^{-16}$
SHAY	-1.2685	0.0010	-1252.07	$< 2 \times 10^{-16}$
SSAN	-1.2774	0.0010	-1256.37	$< 2 \times 10^{-16}$
UCTY	-0.9818	0.0009	-1079.83	$< 2 \times 10^{-16}$

Term	**Estimate**	**Std. Error**	**z value**	**Pr(>|z|)**
WARM	-1.4570	0.0011	-1334.73	$< 2 \times 10^{-16}$
WCRK	-0.5644	0.0008	-716.69	$< 2 \times 10^{-16}$
WDUB	-1.3458	0.0010	-1285.03	$< 2 \times 10^{-16}$
WOAK	-0.3217	0.0007	-439.64	$< 2 \times 10^{-16}$
Hour (reference: hour 0)				
Hour 1	-1.5595	0.0048	-325.39	$< 2 \times 10^{-16}$
Hour 2	-2.2304	0.0210	-106.37	$< 2 \times 10^{-16}$
Hour 3	-3.2591	0.0532	-61.28	$< 2 \times 10^{-16}$
Hour 4	-2.9188	0.0391	-74.73	$< 2 \times 10^{-16}$
Hour 5	-2.9519	0.0119	-248.69	$< 2 \times 10^{-16}$
Hour 6	-0.3479	0.0030	-115.16	$< 2 \times 10^{-16}$
Hour 7	0.3925	0.0025	157.63	$< 2 \times 10^{-16}$
Hour 8	0.7744	0.0023	332.92	$< 2 \times 10^{-16}$
Hour 9	0.9804	0.0023	434.36	$< 2 \times 10^{-16}$
Hour 10	1.1217	0.0022	506.10	$< 2 \times 10^{-16}$
Hour 11	1.2760	0.0022	586.03	$< 2 \times 10^{-16}$
Hour 12	1.3695	0.0022	635.18	$< 2 \times 10^{-16}$
Hour 13	1.3726	0.0022	636.79	$< 2 \times 10^{-16}$
Hour 14	1.3831	0.0022	642.34	$< 2 \times 10^{-16}$
Hour 15	1.4149	0.0021	659.13	$< 2 \times 10^{-16}$
Hour 16	1.4981	0.0021	703.37	$< 2 \times 10^{-16}$
Hour 17	1.5214	0.0021	715.82	$< 2 \times 10^{-16}$
Hour 18	1.4319	0.0021	668.15	$< 2 \times 10^{-16}$
Hour 19	1.1915	0.0022	542.07	$< 2 \times 10^{-16}$
Hour 20	0.8927	0.0023	390.63	$< 2 \times 10^{-16}$
Hour 21	0.7826	0.0023	336.82	$< 2 \times 10^{-16}$
Hour 22	0.6841	0.0024	289.76	$< 2 \times 10^{-16}$
Hour 23	0.4839	0.0024	197.77	$< 2 \times 10^{-16}$
Day (reference: Saturday)				
Sunday	-0.0832	0.0028	-29.99	$< 2 \times 10^{-16}$
Weekday	-0.5872	0.0023	-260.79	$< 2 \times 10^{-16}$
Other predictors				
Holiday (TRUE)	-0.6962	0.0006	-1120.69	$< 2 \times 10^{-16}$
Giants Home (TRUE)	0.0311	0.0002	156.49	$< 2 \times 10^{-16}$
A's Home (TRUE)	0.0076	0.0002	38.77	$< 2 \times 10^{-16}$
Warriors Home (TRUE)	-0.0209	0.0003	-75.24	$< 2 \times 10^{-16}$
Season: Spring	-0.0597	0.0002	-247.49	$< 2 \times 10^{-16}$
Season: Summer	-0.0313	0.0002	-132.15	$< 2 \times 10^{-16}$
Season: Winter	-0.1389	0.0002	-588.60	$< 2 \times 10^{-16}$
Hour × Sunday interaction (reference: hour 0 × Saturday)				
Hour 1 × Sun	0.0043	0.0069	0.62	0.533
Hour 2 × Sun	0.7086	0.0267	26.55	$< 2 \times 10^{-16}$
Hour 3 × Sun	0.8416	0.0729	11.55	$< 2 \times 10^{-16}$
Hour 4 × Sun	-0.9741	0.1124	-8.66	$< 2 \times 10^{-16}$
Hour 5 × Sun	-0.7551	0.0439	-17.21	$< 2 \times 10^{-16}$
Hour 6 × Sun	-2.8156	0.0279	-101.04	$< 2 \times 10^{-16}$
Hour 7 × Sun	-2.0558	0.0062	-332.99	$< 2 \times 10^{-16}$
Hour 8 × Sun	-0.2628	0.0034	-76.36	$< 2 \times 10^{-16}$
Hour 9 × Sun	-0.0918	0.0033	-28.05	$< 2 \times 10^{-16}$

Term	**Estimate**	**Std. Error**	**z value**	**Pr(>|z|)**
Hour 10 × Sun	-0.0517	0.0032	-16.12	$< 2 \times 10^{-16}$
Hour 11 × Sun	-0.0530	0.0031	-16.84	$< 2 \times 10^{-16}$
Hour 12 × Sun	-0.0753	0.0031	-24.15	$< 2 \times 10^{-16}$
Hour 13 × Sun	-0.0958	0.0031	-30.67	$< 2 \times 10^{-16}$
Hour 14 × Sun	-0.1384	0.0031	-44.29	$< 2 \times 10^{-16}$
Hour 15 × Sun	-0.1739	0.0031	-55.71	$< 2 \times 10^{-16}$
Hour 16 × Sun	-0.1615	0.0031	-52.21	$< 2 \times 10^{-16}$
Hour 17 × Sun	-0.1665	0.0031	-53.92	$< 2 \times 10^{-16}$
Hour 18 × Sun	-0.2214	0.0031	-70.85	$< 2 \times 10^{-16}$
Hour 19 × Sun	-0.2373	0.0032	-73.75	$< 2 \times 10^{-16}$
Hour 20 × Sun	-0.2015	0.0033	-60.19	$< 2 \times 10^{-16}$
Hour 21 × Sun	-0.2457	0.0034	-71.78	$< 2 \times 10^{-16}$
Hour 22 × Sun	-0.3708	0.0035	-104.99	$< 2 \times 10^{-16}$
Hour 23 × Sun	-0.4939	0.0037	-132.00	$< 2 \times 10^{-16}$
Hour × Weekday interaction (reference: hour 0 × Saturday)				
Hour 1 × Wkdy	0.1327	0.0056	23.70	$< 2 \times 10^{-16}$
Hour 2 × Wkdy	0.7375	0.0231	31.95	$< 2 \times 10^{-16}$
Hour 3 × Wkdy	1.3586	0.0542	25.08	$< 2 \times 10^{-16}$
Hour 4 × Wkdy	1.0376	0.0393	26.39	$< 2 \times 10^{-16}$
Hour 5 × Wkdy	3.7314	0.0120	312.16	$< 2 \times 10^{-16}$
Hour 6 × Wkdy	2.2338	0.0033	683.17	$< 2 \times 10^{-16}$
Hour 7 × Wkdy	2.2004	0.0028	795.13	$< 2 \times 10^{-16}$
Hour 8 × Wkdy	2.1540	0.0026	823.61	$< 2 \times 10^{-16}$
Hour 9 × Wkdy	1.5990	0.0026	624.56	$< 2 \times 10^{-16}$
Hour 10 × Wkdy	0.9190	0.0025	362.00	$< 2 \times 10^{-16}$
Hour 11 × Wkdy	0.5756	0.0025	229.15	$< 2 \times 10^{-16}$
Hour 12 × Wkdy	0.4765	0.0025	191.05	$< 2 \times 10^{-16}$
Hour 13 × Wkdy	0.5261	0.0025	211.20	$< 2 \times 10^{-16}$
Hour 14 × Wkdy	0.6680	0.0025	268.99	$< 2 \times 10^{-16}$
Hour 15 × Wkdy	0.9285	0.0025	376.10	$< 2 \times 10^{-16}$
Hour 16 × Wkdy	1.2273	0.0024	501.79	$< 2 \times 10^{-16}$
Hour 17 × Wkdy	1.4921	0.0024	612.22	$< 2 \times 10^{-16}$
Hour 18 × Wkdy	1.3096	0.0025	533.01	$< 2 \times 10^{-16}$
Hour 19 × Wkdy	0.9283	0.0025	368.37	$< 2 \times 10^{-16}$
Hour 20 × Wkdy	0.7067	0.0026	269.98	$< 2 \times 10^{-16}$
Hour 21 × Wkdy	0.5787	0.0027	217.13	$< 2 \times 10^{-16}$
Hour 22 × Wkdy	0.5519	0.0027	203.93	$< 2 \times 10^{-16}$
Hour 23 × Wkdy	0.4332	0.0028	154.28	$< 2 \times 10^{-16}$

Table 6: Full coefficient table for the final Poisson GLM. All coefficients are on the log-rate scale. The single non-significant coefficient is Hour 1 × Sunday ($p = 0.533$). Reference levels are: 12TH for destination, hour 0 for hour, and Saturday for day.